

Executive Summary: Automated Cloud Backup & Integrity System

Project Goal: To design, build, and deploy a secure, resilient, and fully automated system to protect critical company data by ensuring all S3-based backups are replicated for disaster recovery and programmatically verified for integrity.

Project Status: Complete and Operational.

Business Problem

As our reliance on cloud storage grows, so does the risk of data loss due to accidental deletion, data corruption, or regional service disruptions. Standard backup procedures often lack automated integrity checks, leaving the business exposed to restoring corrupted or incomplete data. Furthermore, manual or brittle, script-based solutions are inefficient, prone to human error, and do not scale, representing a significant operational risk.

Solution Implemented

I have successfully engineered and deployed a state-of-the-art, event-driven backup validation system entirely on AWS. This system eliminates manual intervention and provides a robust, multi-layered defense for our data assets.

Key capabilities and business benefits include:

- **Enhanced Data Protection & Disaster Recovery:** Every file uploaded is automatically replicated to a secondary, geographically distant AWS region, ensuring business continuity in the event of a regional outage. The integrity of every replicated file is automatically verified, guaranteeing that our backups are always sound.
- **Increased Operational Efficiency:** The entire workflow—from upload to validation to error handling—is 100% automated and serverless. This completely eliminates the manual effort and operational overhead associated with traditional backup management, freeing up engineering resources to focus on value-adding initiatives.
- **Proactive Risk Mitigation:** By integrating a DevSecOps pipeline with automated security scanning (Checkov), I have hardened our infrastructure against common misconfigurations. The system proactively identifies and isolates processing failures into Dead-Letter Queues (DLQs) and sends immediate, actionable alerts to the technical team, transforming our posture from reactive to proactive.

- **Scalability and Cost-Effectiveness:** The serverless, event-driven architecture is inherently scalable, capable of handling a massive influx of data without performance degradation. By leveraging cost-optimization features like S3 Lifecycle Policies, the system automatically archives older data to cheaper storage tiers, ensuring long-term data retention is achieved in the most cost-effective manner.

Technical Foundation

The system was built using a modern, professional-grade technology stack, managed entirely through Infrastructure as Code (Terraform) and deployed via a secure GitHub Actions CI/CD pipeline. This GitOps approach ensures that our infrastructure is repeatable, auditable, and can be managed with the same rigor as our application code.

Outcome & Next Steps

The successful completion of this project has significantly elevated our data protection strategy, reduced operational risk, and improved our overall security posture. The modular and scalable design provides a foundational blueprint that can be extended to protect other critical data sources in the future. The system is now fully operational and requires no day-to-day management, only intervention in the rare event of an unrecoverable processing failure.